Foundations of Natural Language Processing
Lecture 13
Heads, Dependency parsing

Henry S. Thompson

(slides from Alex Lascarides, Henry Thompson, Nathan Schneider and Sharon Goldwater)

5 March 2019

Today we will. . .

- Provide metrics for evaluating a parser

- Return to the problem of PCFGs

- Suggest a fix

- This fix leads to an approach without constituent structure!

<p style="text-align:center; color:red;">Dependency parsing</p>

Evaluating parse accuracy

Compare **gold standard** tree (left) to **parser output** (right):

- Output constituent is counted correct if there is a gold constituent that spans the same sentence positions.

- Harsher measure: also require the constituent labels to match.

- **Pre-terminals** (lexical categories) don't count as constituents.

Evaluating parse accuracy

Compare **gold standard** tree (left) to **parser output** (right):

- **Precision**: (# correct constituents)/(# in parser output) = 3/5

- **Recall**: (# correct constituents)/(# in gold standard) = 3/4

- **F-score**: balances precision/recall: 2pr/(p+r)

Parsing accuracies

F-scores for parsing on WSJ corpus:

- vanilla PCFG: $< 80\%$[1]

- lexicalizing + cat-splitting: 89.5% (Charniak, 2000)

- Best current parsers get about 92%

- Numbers get better if we look at top 5 or top 10

However, results on other corpora and other languages are considerably lower. Definitely not a solved problem!

---

[1]Charniak (1996) reports 81% but using gold POS tags as input.

Summary

- Probabilistic models of syntax can help disambiguation and speed in broad-coverage parsing.

  - by computing the probabilities of each tree or sub-tree as the product of the rules in it, and choosing the best option(s).

- Treebanks provide training data for estimating rule probabilities.

- However, to do well, we need to be clever:

  - Standard categories in the treebank don't capture some important facts about language.
  - By creating more detailed categories, we can encode more information within the PCFG framework.

Recall Problem with Vanilla PCFGs
No lexical dependencies

Replacing one word with another with the same POS will never result in a different parsing decision, even though it should!

- kids saw birds with fish vs.
  kids saw birds with binoculars

- She stood by the door covered in tears vs.
  She stood by the door covered in ivy

- stray cats and dogs vs.
  Siamese cats and dogs

A way to fix PCFGs: lexicalization

Create new categories, this time by adding the **lexical head** of the phrase (note: N level under NPs not shown for brevity):
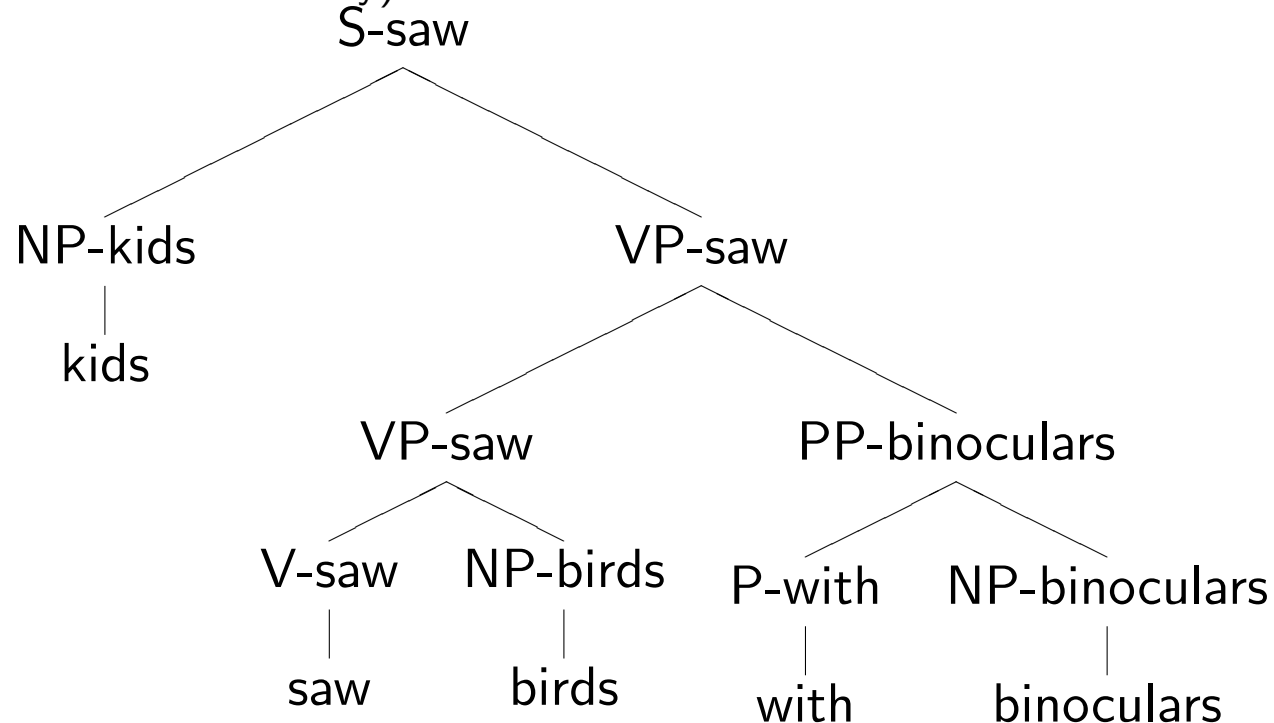
- Now consider:

$\texttt{VP-saw} \rightarrow \texttt{VP-saw PP-fish}$ vs. $\texttt{VP-saw} \rightarrow \texttt{VP-saw PP-binoculars}$
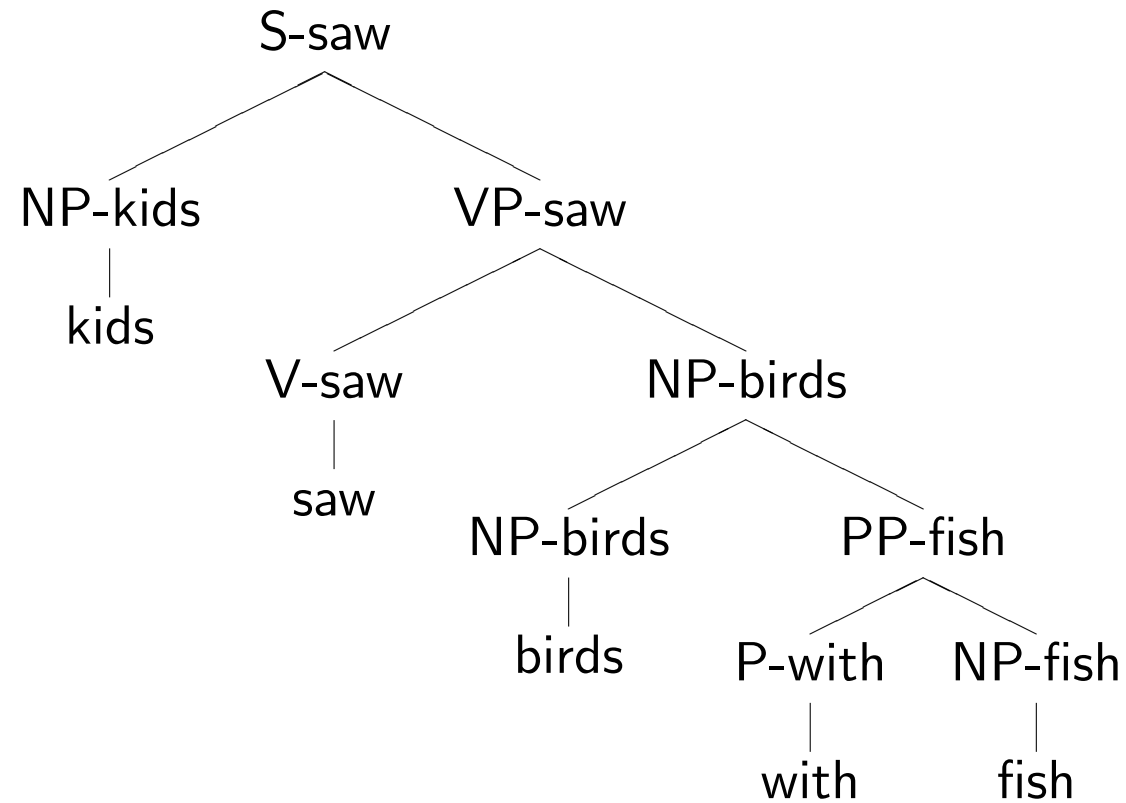
Practical issues

- Identifying the head of every rule is not always straightforward

    - (more on this below)

- All this category-splitting makes the grammar much more **specific** (good!)

- But leads to huge grammar blowup and very sparse data (bad!)

- Lots of effort on how to balance these two issues.

    - Complex smoothing schemes (similar to N-gram interpolation/backoff).
    - More recently, increasing emphasis on automatically learned subcategories.

- But do we really need phrase structure in the first place? Not always!

- Today: Syntax (and parsing) without constituent structure.

Outline

1. Dependencies: what/why

2. Transforming constituency $\rightarrow$ dependency parse

3. Direct dependency parsing

   - Transition-based (shift-reduce)
   - Graph-based

Lexicalized Constituency Parse

```
                          S-saw
                       /         \
                 NP-kids          VP-saw
                   |            /         \
                  kids      V-saw          NP-birds
                             |           /          \
                            saw      NP-birds        PP-fish
                                       |           /        \
                                     birds     P-with      NP-fish
                                                 |            |
                                                with         fish
```

. . . remove the phrasal categories. . .

. . . remove the (duplicated) terminals. . .

```
                    saw
              ┌──────┴──────┐
            kids           saw
                      ┌─────┴─────┐
                     saw         birds
                              ┌───┴────┐
                            birds     fish
                                   ┌───┴───┐
                                 with    fish
```

. . . and collapse chains of duplicates. . .

. . . and collapse chains of duplicates. . .

. . . and collapse chains of duplicates. . .

. . . and collapse chains of duplicates. . .

```
                saw
              /     \
          kids       saw
                    /    \
                 saw    birds
                          |
                        fish
                          |
                        with
```

. . . and collapse chains of duplicates. . .

# . . . and collapse chains of duplicates. . .

# Dependency Parse

Linguists have long observed that the meanings of words within a sentence depend on one another, mostly in *asymmetric*, *binary* relations.

- Though some constructions don't cleanly fit this pattern: e.g., coordination, relative clauses.

Dependency Parse

Equivalently, but showing word order (head → modifier):

Because it is a tree, every word has exactly one parent.

Edge Labels

It is often useful to distinguish different kinds of head $\rightarrow$ modifier **relations**, by labeling edges:

Important relations for English include subject, direct object, determiner, adjective modifier, adverbial modifier, etc. (Different treebanks use somewhat different label sets.)

- How would you identify the subject in a constituency parse?

Dependency Paths

For **information extraction** tasks involving real-world relationships between entities, chains of dependencies can provide good features:

(example from Brendan O'Connor)

Projectivity

- A sentence's dependency parse is said to be **projective** if every subtree (node and all its descendants) occupies a *contiguous span* of the sentence.

- = The dependency parse can be drawn on top of the sentence without any crossing edges.
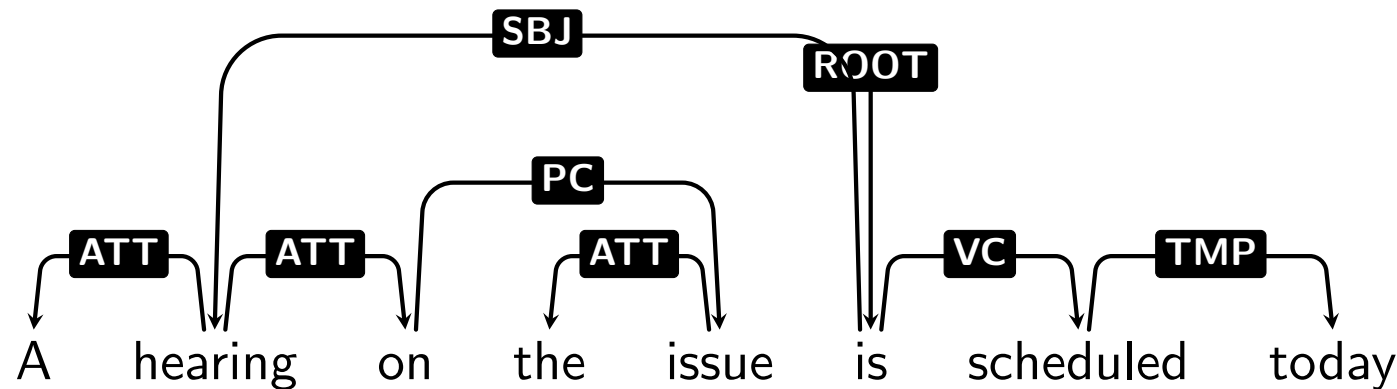
Nonprojectivity

- Other sentences are **nonprojective**:
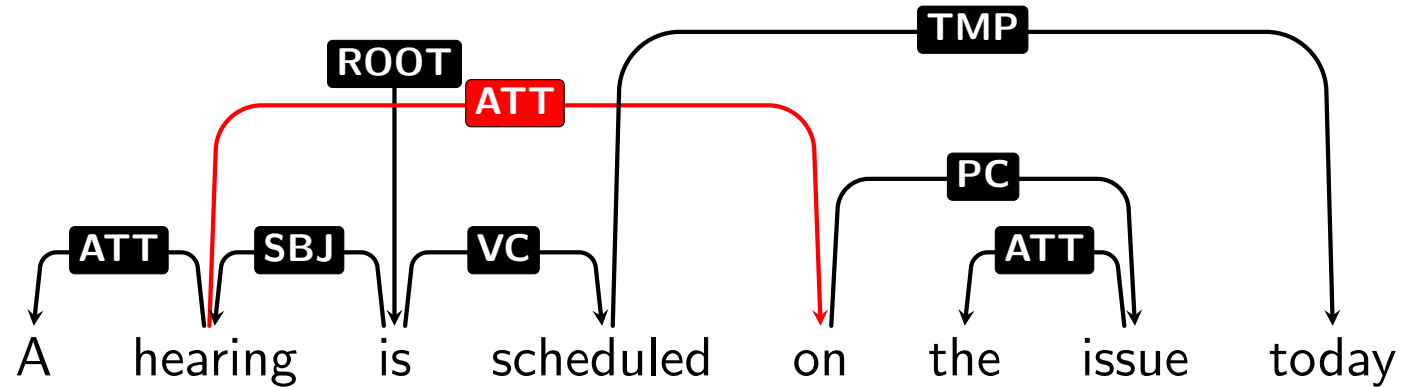
- Nonprojectivity is rare in English, but quite common in many languages.

Outline

1. Dependencies: what/why

2. **Transforming constituency $\rightarrow$ dependency parse**

3. Direct dependency parsing

   - Transition-based (shift-reduce)
   - Graph-based

Constituency Tree → Dependency Tree

We saw how the **lexical head** of the phrase can be used to collapse down to a dependency tree:

```
                            S-saw
                   /                    \
              NP-kids                 VP-saw
                 |              /              \
               kids         VP-saw         PP-binoculars
                          /       \        /           \
                      V-saw    NP-birds  P-with    NP-binoculars
                        |         |        |            |
                      saw       birds    with       binoculars
```
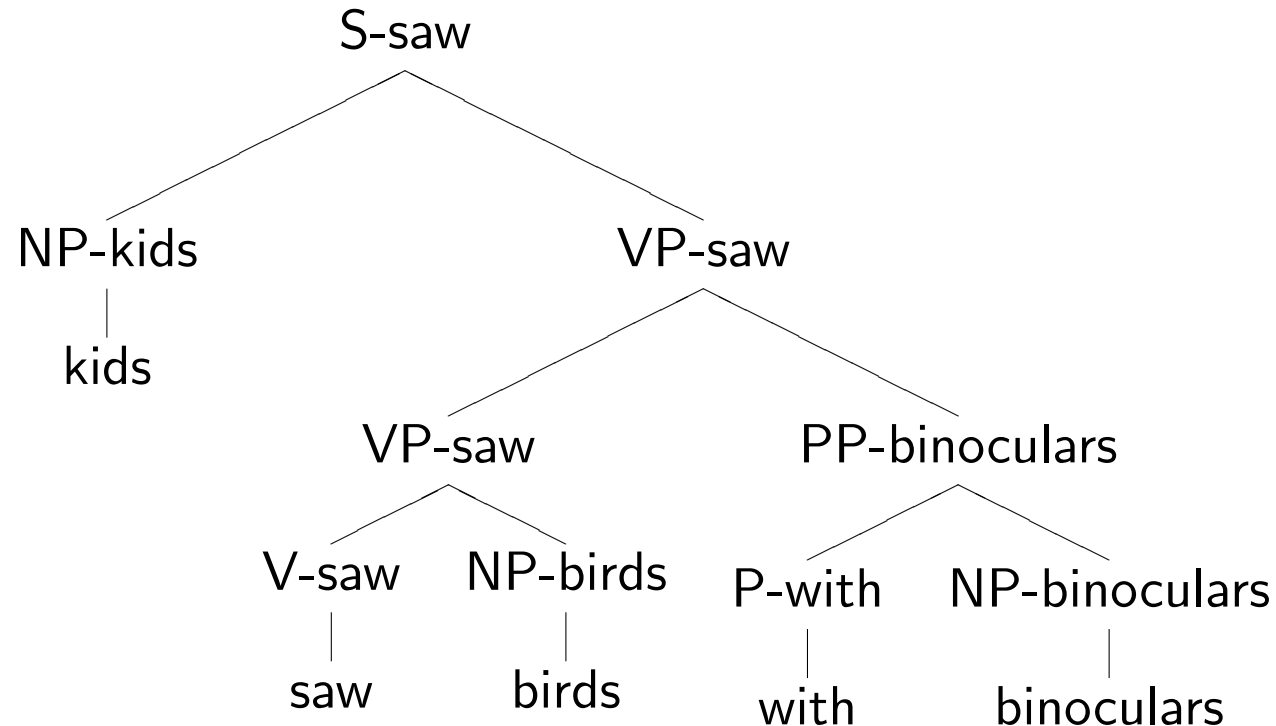
- But how can we find each phrase's head in the first place?

Head Rules

The standard solution is to use **head rules**: for every non-unary (P)CFG production, designate one RHS nonterminal as containing the head. $S \rightarrow NP \; \underline{VP}$, $VP \rightarrow \underline{VP} \; PP$, $PP \rightarrow P \; \underline{NP}$, etc.

- Heuristics to scale this to large grammars: e.g., within an NP, last immediate N child is the head.

Head Rules

Then, propagate heads up the tree:

Head Rules

Then, propagate heads up the tree:

Head Rules

Then, propagate heads up the tree:
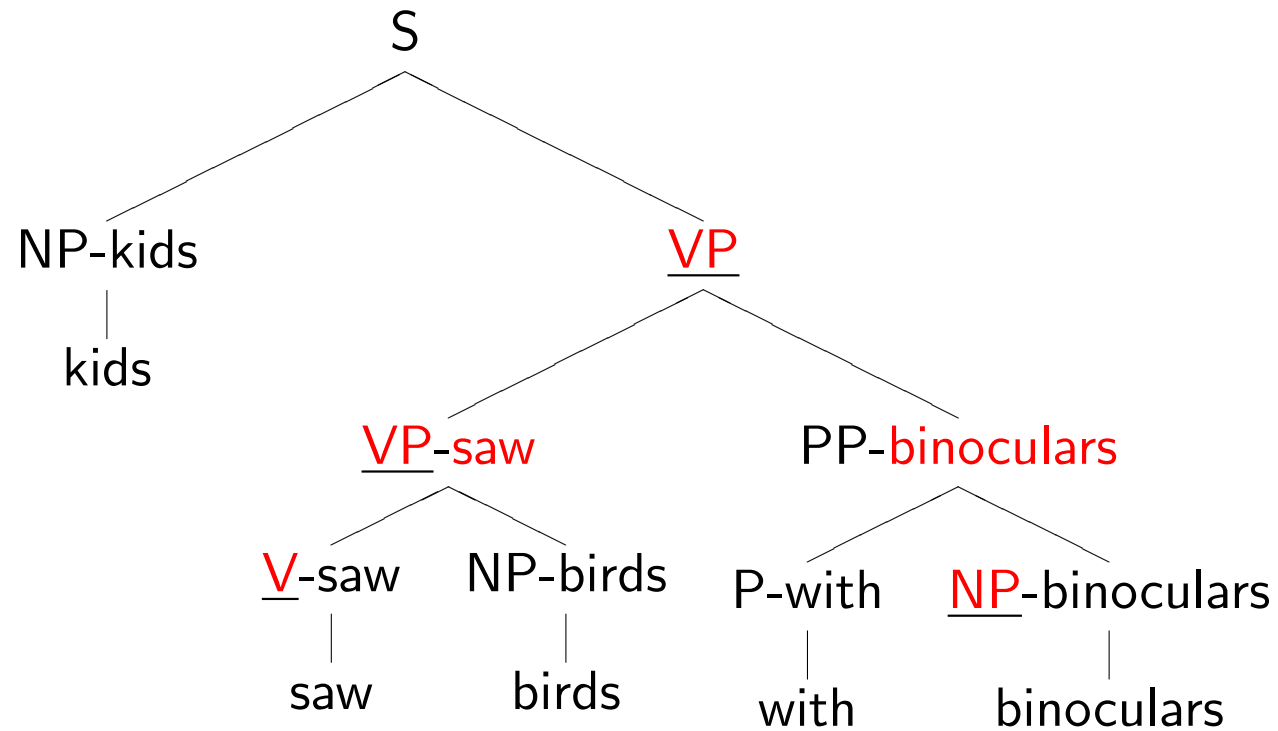
Head Rules

Then, propagate heads up the tree:
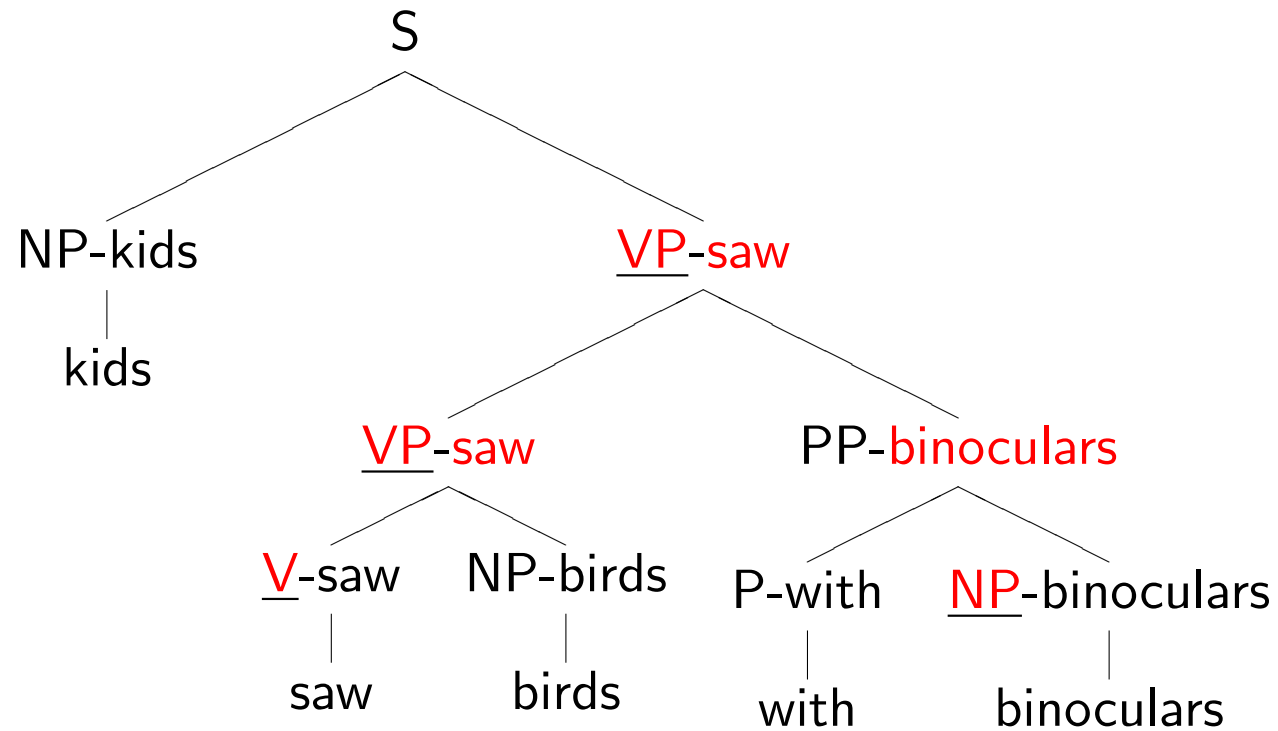
Head Rules

Then, propagate heads up the tree:
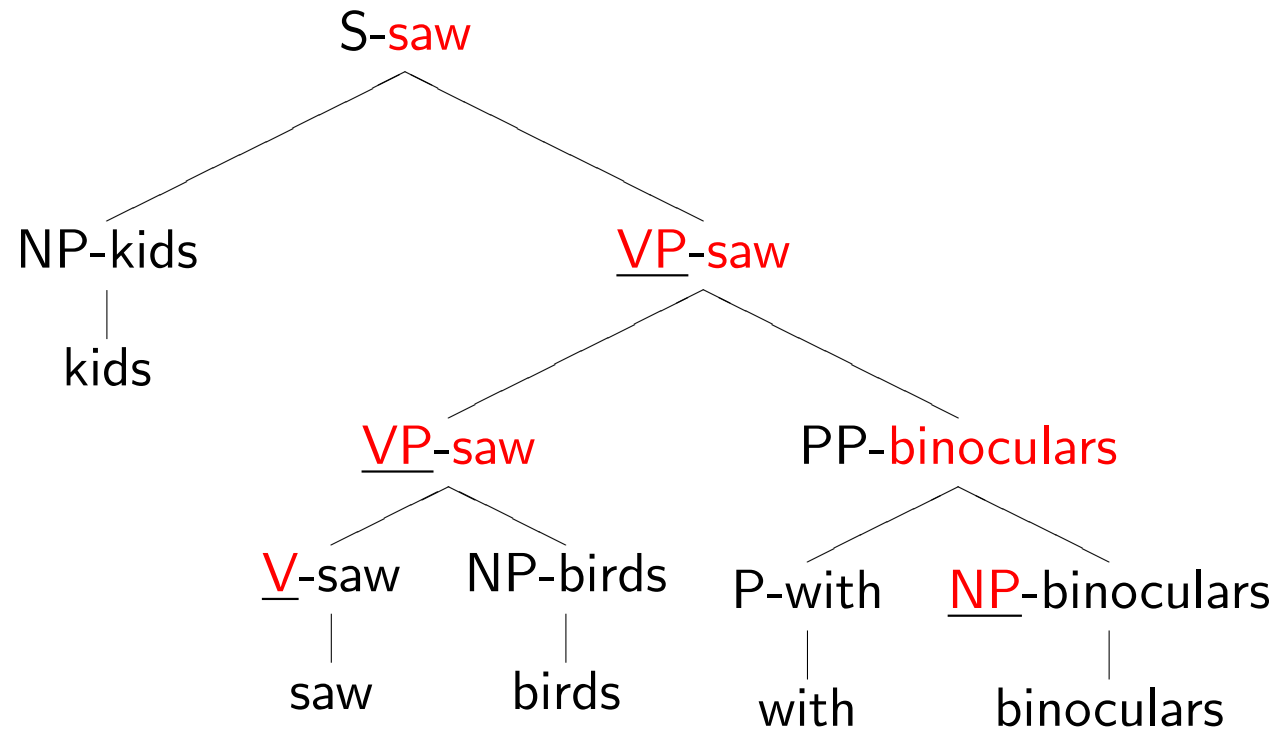
Outline

1. Dependencies: what/why

2. Transforming constituency $\rightarrow$ dependency parse

3. **Direct dependency parsing**

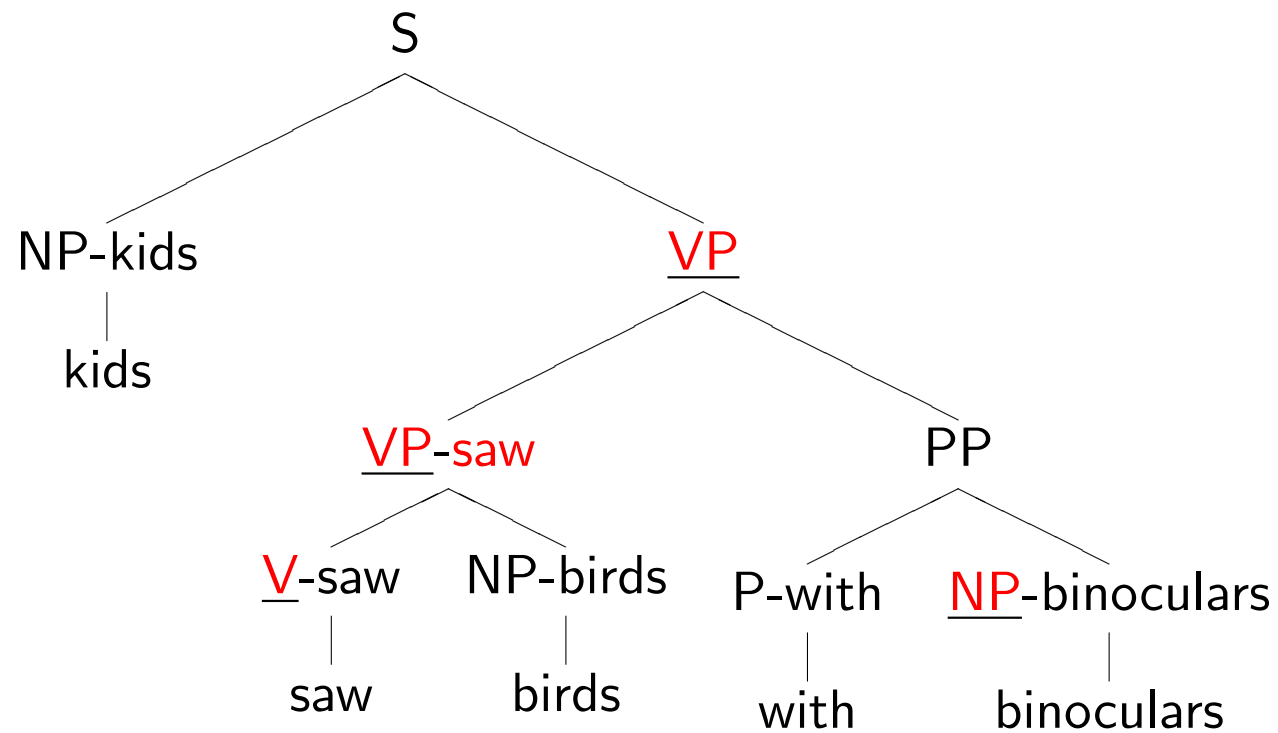   - Transition-based (shift-reduce)
   - Graph-based

Dependency Parsing

Some of the algorithms you have seen for PCFGs can be adapted to dependency parsing.

- **CKY** can be adapted, though efficiency is a concern: obvious approach is $O(Gn^5)$; Eisner algorithm brings it down to $O(Gn^3)$

  - N. Smith's slides explaining the Eisner algorithm: `http://courses.cs.washington.edu/courses/cse517/16wi/slides/an-dep-slides.pdf`

- **Shift-reduce**: more efficient, doesn't even require a grammar!

Transitation-based Parsing: Shift Reduce Parser
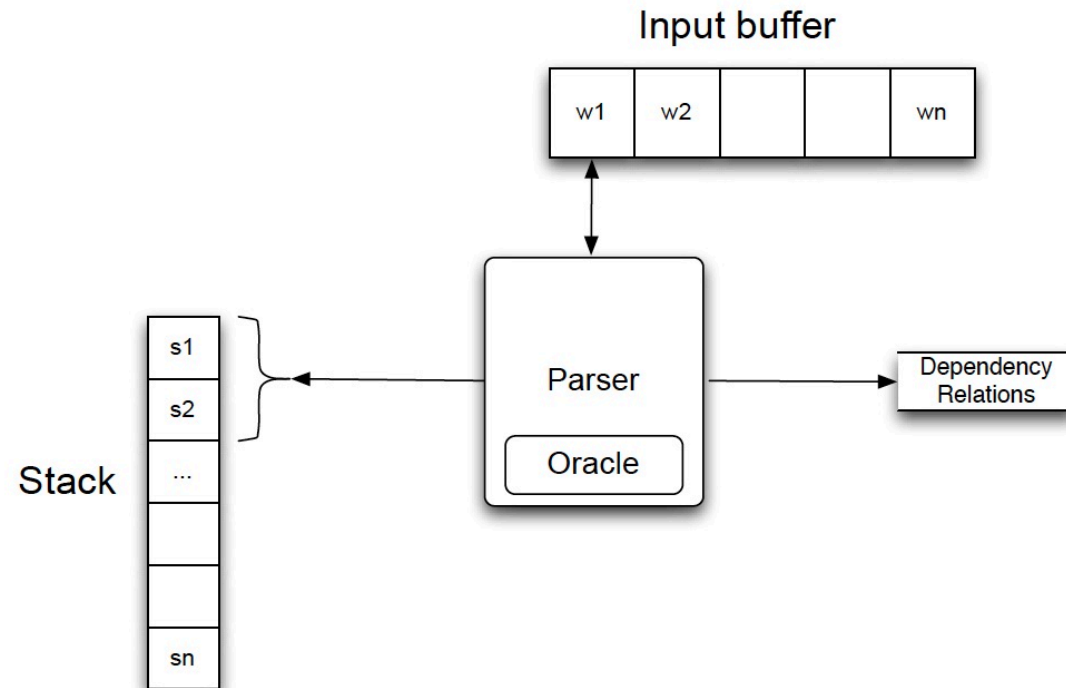
3 possible actions:

LeftArc: Assign head-dependent relation between $s1$ and $s2$; pop $s2$

RightArc: Assign head-dependent relation between $s2$ and $s1$; pop $s1$

Shift: Put $w1$ on top of the stack.

Remember, dependency relation points *from* head *to* dependent

Both **LeftArc** and **RightArc** pop the dependent, leaving the head at the top of the stack

Example

Step	Stack	Word List	Action	Relations
0	[root]	[Kim,saw,Sandy]		
1	[root,Kim]	[saw,Sandy]	Shift	
2	[root,Kim,saw]	[Sandy]	Shift	
3	[root,saw]	[Sandy]	LeftArc	nsubj(saw,Kim)
4	[root,saw,Sandy]	[]	Shift	
5	[root,saw]	[]	RightArc	dobj(saw,Sandy)
6	[root]	[]	RightArc	root→saw

Transition-based Parsing

- Latent structure is just edges between words. Train a **classifier** as the oracle to predict next action (SHIFT, LEFTARC, or RIGHTARC), and proceed left-to-right through the sentence. $O(n)$ time complexity!

- Only finds **projective** trees (without special extensions)

- Pioneering system: Nivre's MALTPARSER

- See `http://spark-public.s3.amazonaws.com/nlp/slides/Parsing-Dependency` `pdf` (Jurafsky & Manning Coursera slides) for details and examples

Graph-based Parsing

- Global algorithm: From the fully connected directed graph of all possible edges, choose the best ones that form a tree.

- **Edge-factored** models: Classifier assigns a nonnegative score to each possible edge; **maximum spanning tree** algorithm finds the spanning tree with highest total score in $O(n^2)$ time.

  - Edge-factored assumption can be relaxed (higher-order models score larger units; more expensive).
  - Unlabeled parse $\rightarrow$ edge-labeling classifier (pipeline).

- Pioneering work: McDonald's MSTPARSER

- Can be formulated as constraint-satisfaction with **integer linear programming** (Martins's TURBOPARSER)

Graph-based vs. Transition-based vs. Conversion-based

- TB: Features in scoring function can look at any part of the stack; no optimality guarantees for search; linear-time; (classically) projective only

- GB: Features in scoring function limited by factorization; optimal search within that model; quadratic-time; no projectivity constraint

- CB: In terms of accuracy, sometimes best to first constituency-parse, then convert to dependencies (e.g., STANFORD PARSER).

  - Slower than direct methods.
  - And, you need a grammar and head rules.

Choosing a Parser: Criteria

- Target representation: constituency or dependency?

- Efficiency? In practice, both runtime and memory use.

- Incrementality: parse the whole sentence at once, or obtain partial left-to-right analyses/expectations?

- Retrainable system?

Choosing a Parser: Performance

SOTA for English constituency parsing (WSJ §23): 91%–92% F_1

Parser	LR	LP	F1	#Toks/s.
Charniak (2000)	89.5	89.9	89.5	–
Klein and Manning (2003)	85.3	86.5	85.9	143
Petrov and Klein (2007)	90.0	90.3	90.1	169
Carreras et al. (2008)	90.7	91.4	91.1	–
Zhu et al. (2013)	90.3	90.6	90.4	1,290
Stanford Shift-Reduce (2014)	89.1	89.1	89.1	655
Hall et al. (2014)	88.4	88.8	88.6	12
This work	89.9	90.4	90.2	957
Charniak and Johnson (2005)[*]	91.2	91.8	91.5	84
Socher et al. (2013)[*]	89.1	89.7	89.4	70
Zhu et al. (2013)[*]	91.1	91.5	91.3	–

Table 3: Results on the English PTB §23. All systems reporting runtimes were run on the same machine. Marked as [*] are reranking and semi-supervised c-parsers.

(Fernández-González and Martins, 2015)

Choosing a Parser: Performance

Constituency parsing in other languages

Parser	Basque	French	German	Hebrew	Hungar.	Korean	Polish	Swedish	Avg.
Berkeley	70.50	**80.38**	78.30	86.96	81.62	71.42	79.23	79.19	78.45
Berkeley Tagged	74.74	79.76	78.28	85.42	85.22	78.56	86.75	80.64	81.17
Hall et al. (2014)	83.39	79.70	78.43	87.18	**88.25**	**80.18**	90.66	82.00	83.72
Crabbé and Seddah (2014)	85.35	79.68	77.15	86.19	87.51	79.35	**91.60**	82.72	83.69
This work	**85.90**	78.75	**78.66**	**88.97**	88.16	79.28	91.20	**82.80**	**84.22**
Björkelund et al. (2014)	88.24	82.53	81.66	89.80	91.72	83.81	90.50	85.50	86.72

(Fernández-González and Martins, 2015)

Choosing a Parser: Performance

SOTA for English dependency parsing (WSJ §23):
93%–94% UAS, 91%–92% LAS (Zhou et al., 2015)

System		UAS	LAS	Speed
baseline greedy parser		91.47	90.43	0.001
Huang and Sagae (2010)		92.10		0.04
Zhang and Nivre (2011)		92.90	91.80	0.03
Choi and McCallum (2013)		92.96	91.93	0.009
Ma et al. (2014)		93.06		
Bohnet and Nivre (2012)†‡		93.67	92.68	0.4
Suzuki et al. (2009)†		93.79		
Koo et al. (2008)†		93.16		
Chen et al. (2014)†		93.77		
beam size				
training	decoding			
100	100	**93.28**	**92.35**	0.07
100	64	93.20	92.27	0.04
100	16	92.40	91.95	0.01

Table 5: Results on WSJ. Speed: sentences per
second. †: semi-supervised learning. ‡: joint
POS-tagging and dependency parsing models.

Labelled **Attachment Score** (LAS) is stricter than *Unlabelled* **Attachment Score** (UAS)

Summary

- While constituency parses give hierarchically nested phrases, dependency parses represent syntax with trees whose edges connect words in the sentence. (No abstract phrase categories like NP.) Edges often labeled with relations like subject.

- Head rules govern how a lexicalized constituency grammar can be extracted from a treebank, and how a constituency parse can be coverted to a dependency parse.

- For English, it is often fastest and most convenient to parse directly to dependencies. Two main paradigms, graph-based and transition-based, with different kinds of models and search algorithms.

- Google "online dependency parser".
  Try out the Stanford parser and SEMAFOR!