Foundations of Natural Language Processing
Lecture 14
Lexical Semantics:
Word senses, relations, and classes

Alex Lascarides
(slides by Lascarides, Schneider, Koehn, Goldwater)

8 March 2019

Meaning

- One grand goal of artificial intelligence is to *understand* what people mean when they talk.

- But how do we know if we succeeded?

Eliza

Joseph Weizenbaum wrote the computer program **Eliza** in 1969 to demonstrate how easily people can be fooled into thinking that machines have some deep understanding.

young woman: Men are all alike.

eliza: In what way?

young woman: They're always bugging us about something specific or other.

eliza: Can you think of a specific example?

young woman: Well, my boyfriend made me come here.

eliza: Your boyfriend made you come here?

young woman: He says I'm depressed much of the time.

eliza: I am sorry to hear you are depressed.

young woman: It's true. I'm unhappy.

eliza: Do you think coming here will help you not to be unhappy?

Online demo: `http://www.manifestation.com/neurotoys/eliza.php3`

What is meaning? What is understanding?

- These are deep philosophical questions

- NLP usually takes a more pragmatic view: can the computer behave as though it understands (in order to do what we want)?

  - Dialogue systems (e.g., Eliza)
  - Machine translation
  - Question answering

- What issues will we face in building such systems?

A Concrete Goal

- We would like to build

  - a machine that answers questions in natural language.
  - may have access to knowledge bases
  - may have access to vast quantities of English text

- Basically, a smarter Google

- This is typically called **Question Answering**

Why is lexical semantics important for building such a system?

Semantics

- To build our QA system we will need to deal with issues in **semantics**, i.e., meaning.

- Lexical semantics: the meanings of individual words (next few lectures)

- Sentential semantics: how word meanings combine (after that)

  - Who did what to whom; when, how, why. . .

- Consider some examples to highlight problems in lexical semantics

Example Question

- Question

  When was Barack Obama born?

- Text available to the machine

  Barack Obama was born on August 4, 1961

- This is easy.

  – just phrase a Google query properly:
    `"Barack Obama was born on *"`
  – syntactic rules that convert questions into statements are straight-forward

Example Question (2)

- Question

  What plants are native to Scotland?

- Text available to the machine

  A new chemical plant was opened in Scotland.

- What is hard?

  – words may have different meanings (**senses**)
  – we need to be able to disambiguate between them

Example Question (3)

- Question

  Where did David Cameron go on vacation?

- Text available to the machine

  David Cameron spent his holiday in Cornwall

- What is hard?

  – words may have the same meaning (**synonyms**)
  – we need to be able to match them

Example Question (4)

- Question

  Which animals love to swim?

- Text available to the machine

  Polar bears love to swim in the freezing waters of the Arctic.

- What is hard?

  – words can refer to a subset (**hyponym**) or superset (**hypernym**) of the concept referred to by another word
  – we need to have database of such **A is-a B** relationships, called an **ontology**

Example Question (5)

- Question

  What is a good way to remove wine stains?

- Text available to the machine

  Salt is a great way to eliminate wine stains

- What is hard?

  – words may be related in other ways, including **similarity** and **gradation**
  – we need to be able to recognize these to give appropriate responses

Example Question (6)

- Question

  Did Poland reduce its carbon emissions since 1989?

- Text available to the machine

  Due to the collapse of the industrial sector after the end of communism in 1989, all countries in Central Europe saw a fall in carbon emissions.

  Poland is a country in Central Europe.

- What is hard?

  – we need to do **inference**
  – a problem for sentential, not lexical, semantics

WordNet

- Some of these problems can be solved with a good ontology, e.g., **WordNet**

- WordNet (English) is a hand-built resource containing 117,000 **synsets**: sets of synonymous words (See `http://wordnet.princeton.edu/`)

- Synsets are connected by relations such as

  - hyponym/hypernym (IS-A: chair-furniture)
  - meronym (PART-WHOLE: leg-chair)
  - antonym (OPPOSITES: good-bad)

- `globalwordnet.org` now lists wordnets in over 50 languages (but variable size/quality/licensing)

Word Sense Ambiguity

- One word form, same category, but more than one sense (**homonyms**):

  I put my money in the *bank.*　　vs.　　He rested at the *bank* of the river.
  I like playing *squash*　　vs.　　I like drinking *squash*

- More generally, words can have multiple (related or unrelated) senses

- Words often exhibit sense ambiguities that fall into (semi-)predictable patterns (**polysemy**): see next slides (from Hugh Rabagliati in PPLS).

Pattern	Participating Senses	Example Sentences
Animal for fur	Mink, chinchilla, rabbit, beaver, raccoon*, alpaca*, crocodile*	The *mink* drank some water / She likes to wear *mink*
Animal/Object for personality	Chicken, sheep, pig, snake, star*, rat*, doll*	The *chicken* drank some water / He is a *chicken*
Animal for meat	Chicken, lamb, fish, shrimp, salmon*, rabbit*, lobster*	The chicken drank some water / The *chicken* is tasty
Artifact for activity	Shower, bath, sauna, baseball,	The *shower* was leaking / The *shower* was relaxing
Body part for object part	Arm, leg, hand, face, back*, head*, foot*, shoulder*, lip*,	John's *arm* was tired / The *arm* was reupholstered
Building for people	Church, factory, school, airplane,	The *church* was built 20 years ago / The *church* sang a song
Complement Coercion	Begin, start, finish, try	John *began* reading the book / John *began* the book
Container for contents	Bottle, can, pot, pan, bowl*, plate*, box*, bucket*	The *bottle* is made of steel / He drank half of the *bottle*
Word for question	Price, weight, speed	The *price* of the coffee was low / John asked the *price* of the coffee

Pattern	Participating Senses	Example Sentences
Figure for Ground	Window, door, gate, goal	The window is broken / The cat walked through the window
Grinding	Apple, chair, fly	The apple was tasty / There is apple all over the table
Instrument for action	Hammer, brush, shovel, tape, lock*, bicycle*, comb*, saw*	The hammer is heavy / She hammered the nail into the wall
Instance of an entity for kind	Tennis, soccer, cat, dog, class*, dinner*, chair*, table*	Tennis was invented in England / Tennis was fun today
Location / Place at location	Bench, land, floor, ground, box*, bottle*, jail*	The bench was made of pine / The coach benched the player
Object for placing at goal	Water, paint, salt, butter, frame*, dress*, oil*	The water is cold / He watered the plant.
Object for taking from source	Milk, dust, weed, peel, pit*, skin*, juice*	The milk tastes good / He milked the cow
Material for artifact	Tin, iron, china, glass, linen*, rubber*, nickel*, fur*	Watch out for the broken glass / He filled the glass with water
Occupation for role in action	Boss, nurse, guard, tutor	My boss is nice / He bossed me around

Pattern	Participating Senses	Example Sentences
Place for an event	Vietnam, Korea, Waterloo, Iraq	It is raining in *Vietnam* / John was shot during *Vietnam*
Place for an institution	White House, Washington, Hollywood, Pentagon, Wall Street*, Supreme Court	The *White House* is being repainted / The *White House* made an announcement
Plant for food or material	Corn, broccoli, coffee, cotton, lettuce*, eggs*, oak*, pine*	The large field of *corn* / The *corn* is delicious
Portioning	Water, beer, jam	She drank some *water* / She bought three *waters*
Publisher for product	Newspaper, magazine, encyclopedia, Wall Street Journal*, New York Times*,	The *newspaper* is badly printed / The *newspaper* fired three employees
Artist for product	Writer, artist, composer, Shakespeare, Dickens*, Mozart*, Picasso*	The *writer* drank a lot of wine / The *writer* is hard to understand
Object for contents	Book, CD, DVD, TV*, magazine*, newspaper*	The heavy, leather- bound *book* / The *book* is funny.
Visual Metaphor	Beam, belt, column, stick, bug*, leaf*	Most of the weight rests on the *beam* / There was a *beam* of light

Some lessons to draw

- Words are typically semantically ambiguous

- There's a lot of regularity (and hence predictability) in the range of senses a word can take

- Those senses also influence the word's syntactic behaviour

- But all regularities admit (arbitrary) exceptions

- Word senses can be **productive**, making a dictionary model (like WordNet) inadequate

- But it's a dominant model in CL these days, and works quite well in lots of cases.

How many senses?

- 5 min. exercise: How many senses does the word interest have?

How many senses?

How many senses?

- How many senses does the word interest have?

  - She pays 3% **interest** on the loan.
  - He showed a lot of **interest** in the painting.
  - Microsoft purchased a controlling **interest** in Google.
  - It is in the national **interest** to invade the Bahamas.
  - I only have your best **interest** in mind.
  - Playing chess is one of my **interests**.
  - Business **interests** lobbied for the legislation.

- Are these seven different senses? Four? Three?

- Also note: distinction between polysemy and homonymy not always clear!

Lexicography requires data

Lumping vs. Splitting

- For any given word, lexicographer faces the choice:

  - **Lump** usages into a small number of senses? or
  - **Split** senses to reflect fine-grained distinctions?

WordNet senses for interest

- S1: a sense of concern with and curiosity about someone or something, Synonym: involvement

- S2: the power of attracting or holding one's interest (because it is unusual or exciting etc.), Synonym: interestingness

- S3: a reason for wanting something done, Synonym: sake

- S4: a fixed charge for borrowing money; usually a percentage of the amount borrowed

- S5: a diversion that occupies one's time and thoughts (usually pleasantly), Synonyms: pastime, pursuit

- S6: a right or legal share of something; a financial involvement with something, Synonym: stake

- S7: (usually plural) a social group whose members control some field of activity and who have common aims, Synonym: interest group

Synsets and Relations in WordNet

- **Synsets** ("synonym sets", effectively senses) are the basic unit of organization in WordNet.

  – Each synset is specific to nouns (`.n`), verbs (`.v`), adjectives (`.a`, `.s`), or adverbs (`.r`).
  – Synonymous words belong to the same synset: car^1 (`car.n.01`) $=$ {car,auto,automobile}.
  – Polysemous words belong to multiple synsets: car^1 vs. $\text{car}^4 =$ {car,elevator car}. Numbered roughly in descending order of frequency.

- Synsets are organized into a **network** by several kinds of relations, including:

  – **Hypernymy** (Is-A): hyponym {ambulance} is a kind of hypernym car^1
  – **Meronymy** (Part-Whole): meronym {air bag} is a part of holonym car^1

Visualizing WordNet
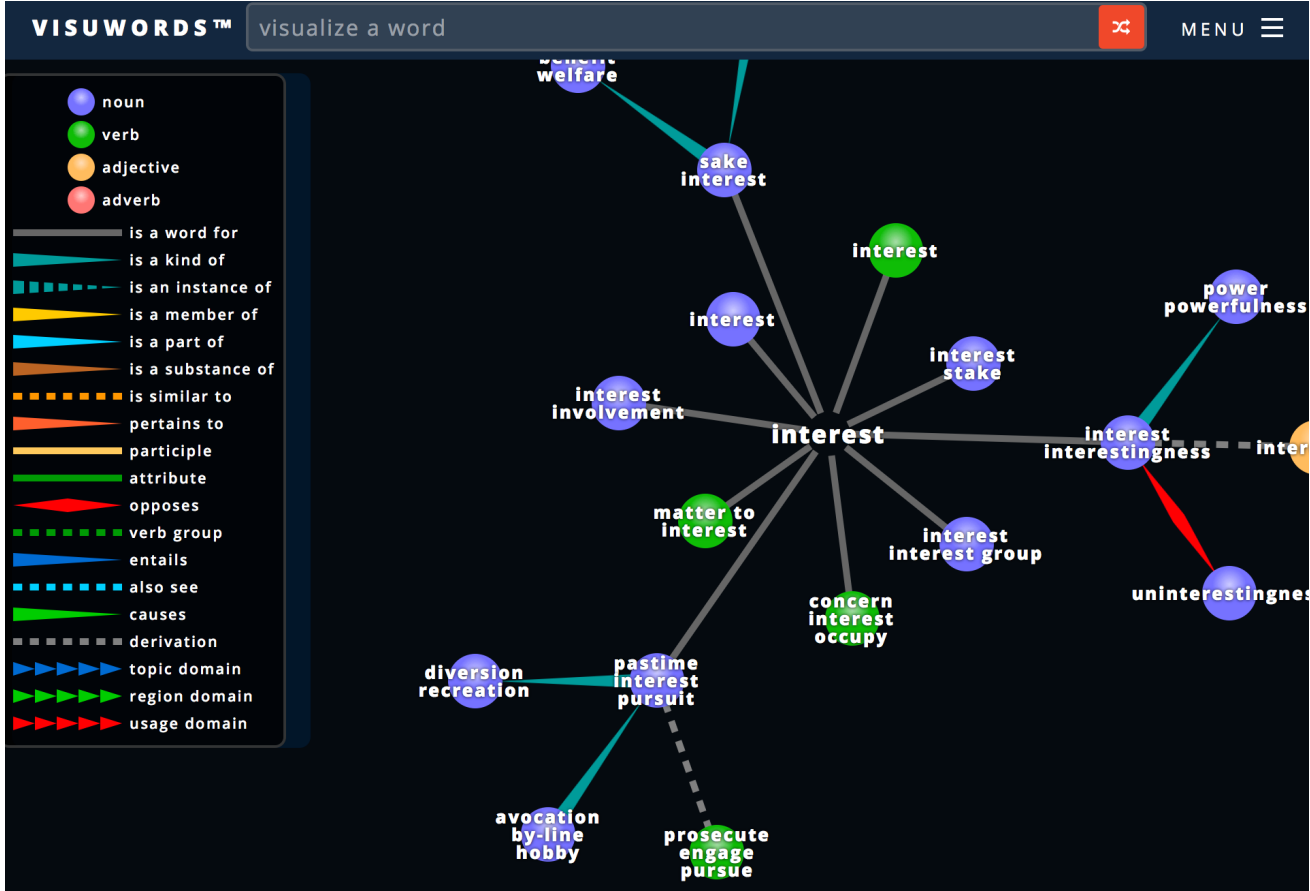
Using WordNet

- NLTK provides an excellent API for looking things up in WordNet:

```
>>> from nltk.corpus import wordnet as wn
>>> wn.synsets('car')
[Synset('car.n.01'), Synset('car.n.02'), Synset('car.n.03'),
Synset('car.n.04'), Synset('cable_car.n.01')]
>>> wn.synset('car.n.01').definition()
u'a motor vehicle with four wheels; usually propelled by an
internal combustion engine'
>>> wn.synset('car.n.01').hypernyms()
[Synset('motor_vehicle.n.01')]
```

- (WordNet uses an obscure custom file format, so reading the files directly is not recommended!)

Coverage in WordNet

- Online stats:

  - 155k unique strings, 118k unique synsets, 207k pairs
  - nouns have an average 1.24 senses (2.79 if exluding monosemous words)
  - verbs have an average 2.17 senses (3.57 if exluding monosemous words)

- Too fine-grained?

- WordNet is a snapshot of the English lexicon, but by no means complete.

  - E.g., consider **multiword expressions** (including noncompositional expressions, idioms): hot dog, take place, carry out, kick the bucket are in WordNet, but not take a break, stress out, pay attention
  - Neologisms: hoodie, facepalm
  - Names: Microsoft
  - Predictable but novel uses of known words: Badger is a delicacy in Mongolia.

Different sense = different translation

- Another way to define senses: if occurrences of the word have different translations, these indicate different sense

- Example interest translated into German

  - Zins: financial charge paid for load (WordNet sense 4)
  - Anteil: stake in a company (WordNet sense 6)
  - Interesse: all other senses

- Other examples might have distinct words in English but a polysemous word in German.

Word sense disambiguation (WSD)

- For many applications, we would like to disambiguate senses

  - we may be only interested in one sense
  - searching for chemical plant on the web, we do not want to know about chemicals in bananas

- Task: Given a sense ambiguous word, find the sense in a given *context*

- Popular topic, data driven methods perform well

WSD as classification

- Given a word token in context, which sense (class) does it belong to?

- We can train a supervised classifier, assuming sense-labeled training data:

  - She pays 3% **interest/INTEREST-MONEY** on the loan.
  - He showed a lot of **interest/INTEREST-CURIOSITY** in the painting.
  - Playing chess is one of my **interests/INTEREST-HOBBY**.

- **SensEval** and later **SemEval** competitions provide such data

  - held every 1-3 years since 1998
  - provide annotated corpora in many languages for WSD and other semantic tasks

What kind of classifier?

Lots of options available:

- Naïve Bayes, MaxEnt (see Lecture 7)

- Decision lists (see J&M, 20.2.2)

- Decision trees (see any ML textbook)

Naïve Bayes for WSD

$$
\begin{aligned}
\hat{s} &= \arg\max_{s \in S} P(s|\vec{f}) & \\
&= \arg\max_{s \in S} \frac{P(\vec{f}|s)P(s)}{P(\vec{f})} & \text{Bayes} \\
&= \arg\max_{s \in S} P(\vec{f}|s)P(s) & P(\vec{f}) \text{ is fixed} \\
&\approx \arg\max_{s \in S} P(s) \prod_{j=1}^{n} P(f_j|s) & \text{cond. independence}
\end{aligned}
$$

- Naïve Bayes requires estimates of:

  – The prior probability of each class (sense)
  – The probability of each feature given each class

- These can be estimated from the training data.

- But what features to use? (Same question for other classifiers!)

Simple features

- Directly neighboring words (and/or their lemmas)

  - **interest** paid
  - rising **interest**
  - lifelong **interest**
  - **interest** rate
  - **interest** piqued

- Any content words in a 50 word window

  - pastime
  - financial
  - lobbied
  - pursued

More features

- Syntactically related words

- Syntactic role in sense

- Topic of the text

- Part-of-speech tag, surrounding part-of-speech tags

Of course, with NB we have the usual problem with correlated features. MaxEnt doesn't assume they are independent.

Evaluation

- Extrinsic: test as part of IR, QA, or MT system

- Intrinsic: evaluate classification accuracy or precision/recall against gold-standard senses

- Baseline: choose the most frequent sense (sometimes hard to beat)

Issues with WSD

- Not always clear how fine-grained the gold-standard should be

- Difficult/expensive to annotate corpora with fine-grained senses

- Classifiers must be trained separately for each word

  - Hard to learn anything for infrequent or unseen words
  - Requires new annotations for each new word
  - Motivates unsupervised and semi-supervised methods (see J&M 20.5, 20.10)

Semantic Classes

- Other approaches, such as **named entity recognition** and **supersense tagging**, define coarse-grained semantic categories like PERSON, LOCATION, ARTIFACT.

- Like senses, can disambiguate: APPLE as ORGANIZATION vs. FOOD.

- Unlike senses, which are *refinements* of particular words, classes are typically larger groupings.

- Unlike senses, classes can be applied to words/names not listed in a lexicon.

Named Entity Recognition

- Recognizing and classifying **proper names** in text is important for many applications. A kind of **information extraction**.

- Different datasets/named entity recognizers use different inventories of classes.
  - Smaller: PERSON, ORGANIZATION, LOCATION, MISCELLANEOUS
  - Larger: sometimes also PRODUCT, WORK_OF_ART, HISTORICAL_EVENT, etc., as well as numeric value types (TIME, MONEY, etc.)

- NER systems typically use some form of feature-based sequence tagging, with features like capitalization being important.

- Lists of known names called **gazetteers** are also important.

Supersenses in WordNet

N:TOPS	N:OBJECT	V:COGNITION
N:ACT	N:PERSON	V:COMMUNICATION
N:ANIMAL	N:PHENOMENON	V:COMPETITION
N:ARTIFACT	N:PLANT	V:CONSUMPTION
N:ATTRIBUTE	N:POSSESSION	V:CONTACT
N:BODY	N:PROCESS	V:CREATION
N:COGNITION	N:QUANTITY	V:EMOTION
N:COMMUNICATION	N:RELATION	V:MOTION
N:EVENT	N:SHAPE	V:PERCEPTION
N:FEELING	N:STATE	V:POSSESSION
N:FOOD	N:SUBSTANCE	V:SOCIAL
N:GROUP	N:TIME	V:STATIVE
N:LOCATION	V:BODY	V:WEATHER
N:MOTIVE	V:CHANGE	

- The **supersense tagging** goes beyond NER to cover all nouns and verbs.

Summary (1)

- In order to support technologies like question answering, we need ways to reason computationally about **meaning**. **Lexical semantics** addresses meaning at the word level.

  – Words can be ambiguous, sometimes with related meanings (**polysemy**), and other times with unrelated meanings (**homonymy**).
  – Different words can mean the same thing (**synonymy**).

- Computational lexical databases, notably WordNet, organize words in terms of their meanings.

  – **Synsets** and relations between them such as hypernymy and meronymy.

Summary (2)

- **Word sense disambiguation** is the task of choosing the right sense for the context.

  - Classification with contextual features
  - Relying on dictionary senses has limitations in granularity and coverage

- **Semantic classes**, as in NER and supersense tagging, are a coarser-grained representation for semantic disambiguation and generalization.

Next Lecture: Distributional lexical semantics

- What can we learn about a word's meaning from "the company it keeps"?

- What do we do if our thesaurus is incomplete?

- Distributional lexical semantics is about learning word meaning from the contexts in which words appear