

---

# Foundations of Natural Language Processing

## Lecture 17

### Discourse Coherence

Alex Lascarides

19 March 2019



# Making sense of actions



# Changing our minds



# Observing Action

- We assume action choice isn't arbitrary; choice is informed by the context
- So we infer more than we see.
- And may change these inferences as we see more.

# Coherence in Discourse: Making sense of verbal actions



It's a beautiful night.  
We're looking for something dumb to do.  
Hey baby, I think I wanna marry you.

# Questions

Coherence and Content

**Representation:** How should discourse coherence be represented formally and computationally?

**Construction:** What inference processes, and what knowledge sources, are used when identifying coherence relations?

# Outline

- Motivation for Discourse Coherence
- Representing Discourse Coherence
- Inferring Discourses Coherence

# Pronouns

From Hobbs (1985)

John can open Bill's safe.  
He knows the combination



# Pronouns

From Hobbs (1985)

John can open Bill's safe.

John ~~He~~ knows the combination.

- If “He” is John: Explanation (“because”).

# Pronouns

From Hobbs (1985)

John can open Bill's safe.

**Bill** ~~He~~ knows the combination.

- If “He” is John: **Explanation** (“because”).  
If “He” is Bill: at best we infer **Continuation** (“and”) with a very vague topic.

# Pronouns

From Hobbs (1985)

John can open Bill's safe.

He should change the combination.

# Pronouns

From Hobbs (1985)

John can open Bill's safe.

**Bill** ~~He~~ should change the combination.

- If “He” is Bill: **Result** (“so”)

# Pronouns

From Hobbs (1985)

John can open Bill's safe.

**John** ~~He~~ should change the combination.

- If “He” is Bill: **Result** (“so”)  
If “He” is John: a ‘weaker’ **Result**?
- Subjects are more likely antecedents, but not here...

## Pronouns and Coherence

- Pronouns interpreted in a way that maximises coherence, even if this conflicts with predictions from other knowledge sources!

# Coherence and Time

Max fell. John helped him up.

Max fell. John pushed him.

# Coherence and Time

John hit Max on the back of his neck.

Max fell. John pushed him.

Max rolled over the edge of the cliff.

# Word Meaning

A: Did you buy the apartment?

B: Yes, but we rented it./ No, but we rented it.



# Bridging

John took an engine from Avon to Dansville.  
He picked up a boxcar./He also took a boxcar.

# Discourse Coherence and Implicit Agreement

From Sacks *et al.* (1974):

- (1)
  - a. M (to K and S): Karen 'n' I're having a fight,
  - b. M (to K and S): after she went out with Keith and not me.
  - c. K (to M and S): Wul Mark, you never asked me out.

# Discourse Coherence and Dishonesty

Example from Solan and Tiersma (2005)

- (2)    a.    P: Do you have any bank accounts in Swiss banks, Mr. Bronston?
  - b.    B: No, sir.
  - c.    P: Have you ever?
  - d.    B: The company had an account there for about six months, in Zurich.
- (2)d interpreted as an indirect answer, implying *no*...

# Discourse Coherence and Dishonesty

Example from Solan and Tiersma (2005)

- (2)
- a. P: Do you have any bank accounts in Swiss banks, Mr. Bronston?
  - b. B: No, sir.
  - c. P: Have you ever?
  - d. B: The company had an account there for about six months, in Zurich.

- (2)d interpreted as an indirect answer, implying *no*...
- ... *even if* you know it conflicts with Bronston's beliefs.


# Discourse Coherence and Dishonesty

Example from Solan and Tiersma (2005)

- (2)
- a. P: Do you have any bank accounts in Swiss banks, Mr. Bronston?
  - b. B: No, sir.
  - c. P: Have you ever?
  - d. B: The company had an account there for about six months, in Zurich.

- (2)d interpreted as an indirect answer, implying *no*...
- ... *even if* you know it conflicts with Bronston's beliefs.
- Literally true, but negative answer false.
- Supreme court overruled conviction for perjury.
- Different ruling probable if Bronston had said "only".


# Gesture



Now one thing you could do is totally audiotape hours and hours...

*... so that you get a large amount of data that you can think of as laid out on a time line.*


# Gesture



And exhaustively go through and make sure that you really pick up all the speech errors

*... by individually analysing each acoustic unit along the timeline of your data.*

# Gesture



Allow two different coders to go through it...

*... and moreover get them to work independently and reconcile their activities.*



## Meaning of Multimodal Communicative Actions

Coherence relations connect speech and gesture and sequences of gestures.

- speech **so that** gesture  
speech **by** gesture  
speech **and moreover** gesture

# SDRT: The logical form (LF) of monologue

LF consists of:

1. Set  $A$  of **labels**  $\pi_1, \pi_2, \dots$   
(each label stands for a segment of discourse)
2. A **mapping**  $\mathcal{F}$  from each label to a formula representing its content.
3. Vocabulary includes coherence relations; e.g., *Elaboration*( $\pi_1, \pi_2$ ).

## LFs and Coherence

Coherent discourse is a single segment of rhetorically connected subsegments. More formally:

- The **partial order** over  $A$  induced by  $\mathcal{F}$  has a **unique root**.

# An Example

$\pi_1$ : John can open Bill's safe.

$\pi_2$ : He knows the combination.

$\pi_0$  : *Explanation*( $\pi_1, \pi_2$ )

$\pi_1$  :  $\iota x(\text{safe}(x) \ \& \ \text{possess}(x, \text{bill}) \ \& \ \text{can}(\text{open}(e_1, \text{john}, x)))$

$\pi_2$  :  $\iota y(\text{combination}(y) \ \& \ \text{of}(y, x) \ \& \ \text{knows}(\text{john}, y))$

- Bits in **red** are specific values that go beyond content that's revealed by linguistic form.
- They are inferred via **commonsense reasoning** that's used to construct a **maximally coherent** interpretation.

# SDRT: Logical form of dialogue Lascarides and Asher (2009)

- LF tracks all current **public commitments** for each agent, including commitments to **coherence relations**.

- (1)
- a. M (to K and S): Karen 'n' I're having a fight,
  - b. M (to K and S): after she went out with Keith and not me.
  - c. K (to M and S): Wul Mark, you never asked me out.

Turn	M	K
1	$\pi_{1M} : \textit{Explanation}(a, b)$	$\emptyset$
2	$\pi_{1M} : \textit{Explanation}(a, b)$	$\pi_{2K} : \textit{Explanation}(a, b) \wedge \textit{Explanation}(b, c)$

- (2)
- a. P: Do you have any bank accounts in Swiss banks?
  - b. B: No, sir.
  - c. P: Have you ever?
  - d. B: The company had an account there for 6 months.

Turn	Prosecutor	Bronston
1	$a : \mathcal{F}(a)$	$\emptyset$
2	$a : \mathcal{F}(a)$	$\pi_{2B} : \text{Answer}(a, b)$
3	$\pi_{3P} : \text{Continuation}(a, c)$	$\pi_{2B} : \text{Answer}(a, b)$
4	$\pi_{3P} : \text{Continuation}(a, c)$	$\pi_{4B} : \text{Answer}(a, b) \wedge \text{Continuation}(a, c) \wedge \text{Indirect-Answer}(c, d)$

1. **Plausible Deniability**: Must test rigorously whether it's safe to treat the implied answer as a matter of public record.

- (2)
- a. P: Do you have any bank accounts in Swiss banks?
  - b. B: No, sir.
  - c. P: Have you ever?
  - d. B: The company had an account there for 6 months.

Turn	Prosecutor	Bronston
1	$a : \mathcal{F}(a)$	$\emptyset$
2	$a : \mathcal{F}(a)$	$\pi_{2B} : \text{Answer}(a, b)$
3	$\pi_{3P} : \text{Continuation}(a, c)$	$\pi_{2B} : \text{Answer}(a, b)$
4	$\pi_{3P} : \text{Continuation}(a, c)$	$\pi_{4B} : \text{Answer}(a, b) \wedge \text{Continuation}(b, d)$

1. **Plausible Deniability**: Must test rigorously whether it's safe to treat the implied answer as a matter of public record.
2. **Neologism proof equilibria**: distinguishes (2)d vs. "only".

# Symbolic approaches to constructing LF

- Draw on rich information sources:
  - linguistic content, world knowledge, mental states. . .
- Deploy reasoning that supports inference with partial information. Unlike classical logic, this requires **consistency tests**.
- Typically, construct LF and evaluate it **in the same logic**, making constructing LF **undecidable**.

# Further Problem

- Like any knowledge rich approach involving hand-crafted rules, this is only feasible for very small domains.
- Ideally, we would like to **learn** a discourse parser automatically from corpus data.
- But there's a lack of corpora annotated with discourse structure.
  - RSTbank, Graphbank, Annodis, STAC are relatively small.
  - Discourse Penn Treebank is relatively large but not annotated with complete discourse structure.
  - Groningen Parallel Meaning Bank: full discourse structure (SDRSs) and getting bigger all the time.



# Supervised Learning for SDRT

Training on 100 dialogues

Baldrige and Lascarides (2005)

Parser based on Collins' parsing model:

- 72% f-score on segmentation (baseline: 53.3%)
- 48% f-score on segmentation and coherence relations (baseline: 7.4%)
- Doesn't attempt to estimate LFs of clauses.

Training on Groningen Meaning Bank

Liu and Lapata (2018)

Neural semantic parser, RNN computes structure first, fills in arguments later:

- 77% f-score on segmentation, coherence relations *and* LFs of clauses
- State of the Art!

- Coherence relations can be overtly signalled:
  - *because* signals EXPLANATION; *but* signals CONTRAST
- So produce a training set *automatically*:
  - Max fell because John pushed him  
⇒  
EXPLANATION(*Max fell, John pushed him*).

# Results of Best Model

- Test examples originally had a cue phrase: 60.9%.
- Test examples originally had **no cue phrase**: 25.8%
- Train on 1K manually labelled examples: 40.3%.
- Combined training set of manual and automatically labelled examples doesn't improve accuracy.

**So you're better off manually labelling a small set of examples!**

# Results of Best Model

- Test examples originally had a cue phrase: 60.9%.
- Test examples originally had **no cue phrase**: 25.8%
- Train on 1K manually labelled examples: 40.3%.
- Combined training set of manual and automatically labelled examples doesn't improve accuracy.

**So you're better off manually labelling a small set of examples!**

Why?

## *Contrast to Elaboration*

**Although** the electronics industry has changed greatly, possibly the greatest change is that very little component level manufacture is done in this country.

# Conclusion

- Interpretation governed by discourse coherence:
  - Constrains what can be said next
  - Augments meaning revealed by linguistic form.
- Computing logical form should be decidable; modularity is key to this.
- Data-driven approaches are a major challenge.
- Linking rich models of discourse semantics to models of human behaviour and decision making is also a major challenge, but essential for tackling dialogues where the agents' goals conflict.